

LinkedIn Post Generator



Turn insights from Lenny's Podcast into compelling LinkedIn posts that spark conversations, ***without AI slop*** and grow your professional audience.

Zahra
Parisa
Jay
Mudit

FROM TOPIC TO LINKEDIN POST IN MINS

*Extract Insights, Drive
Engagement.*



1. Enter a topic
2. Tell us the language, word count
3. Pick structure / template of the topic
4. Optionally, pick tone, style, goal, target audience, call to action
5. lennybuddies extracts the sharpest insights, frameworks, and quotes — then transforms them into scroll-stopping LinkedIn posts complete with strong hooks, structured takeaways, and relevant hashtags ready to publish.

TOPIC, EXTRACT, PUBLISH

lennybuddies



ARCHITECTURE



•

KNOWLEDGE SOURCES

1. Lenny's curated newsletters-10 editions
2. Lenny's full podcast transcripts from 50 episodes

For Summarisation Layer

1. Python libraries-
 - a. BART X
 - b. T5 X
2. Chunking X
3. Use a Better Pre-trained Summarizer
 - a. Open source models-ollama, mistral X
 - b. Closed source models gpt-4o-mini ✓

For Filtering Layer

1. TF-IDF + cosine similarity (scikit-learn) ✓

Learnings:

1. No-RAG decision helped us to stay lightweight
2. Experimenting with various options gave us a flavor of how to separate good from bad before scaling

RAG DECISION

Why Non-RAG?

- **Corpus size:** ~60 (10 newsletters + 50 podcast episodes) — small enough to reason about without a vector database
- **Complexity vs. scope:** a full vector RAG stack (embeddings + FAISS/Chroma) was unnecessary overhead for this timeframe
- **Approach:** TF-IDF + cosine similarity (scikit-learn) locally scores every document against the user's topic — no embedding API, no vector store, no added latency

Why this isn't RAG

- TF-IDF/cosine similarity is classical sparse keyword matching, not the dense-embedding retrieval most people mean by "RAG"
- We treat it as a lightweight relevance filter inside a non-RAG context-injection design
- Grounding + traceability still hold: every post cites its source article, no hallucinated content

When we'd revisit this

- If the corpus grew substantially larger, or if TF-IDF started missing semantically related content that doesn't share exact keywords with the topic

TECHNICAL IMPLEMENTATION

Tech Stack

- OpenAI GPT-4o-mini
- TF-IDF + cosine similarity (scikit-learn)
- Python backend
- Gradio front-end UI
- GitHub for version control
- VS code + Codex/Claude coding assistant
- Trello for WBS and task tracking

Prompt Design & Agents

- System prompts tuned for Lenny-style tone: actionable, concise, product-focused
- Chain-of-thought prompting for structured LinkedIn post generation
- Review draft content for AI slop

WHAT NEXT


Scope expansion:

- Templates library — Users can add their **own templates** for specific types of posts
- Fan/Influencer-inspired posts — Users can add their **favourite writers/influencers to take inspiration** for tone, language, brevity, style, etc
- Prompt sensitivity — minor wording changes in topics caused **major tone shifts**; requires extensive iteration
- Prompt consistency — tone and voice changes, when changed, **did not have a huge difference** in the post, requires more prompt testing
- Evaluation difficulty — measuring "good" LinkedIn post quality is subjective **without clear metrics** so will have to build in some **benchmarks**
- Logging for relevance, popularity — log which of the retrieved pieces actually influenced accepted posts — over time that tells us which newsletters and podcasts are **earning their place in the knowledge base**
- Publish on LinkedIn — Post generated posts directly on LinkedIn

lennybuddies

START CREATING POSTS NOW

*Generate Your First Post.
Make Your Voice Heard.*



*Thank
You*